

First Impressions. When I first started the SVM image classification exercise, I was a bit nervous but mostly excited. I've worked with SVM before (ITAI 1370), but not for image classification. I wasn't sure how well it would handle images, but I was curious to see it in action. I tried to do it without the given subset of the exercise, but it was taking too long.

Learning Process. Working through the notebook was straightforward for the most part. Concepts like converting images to grayscale and normalizing them made sense, and the code examples helped a lot. But the idea of applying SVM, which I associate more with structured data, to image classification was tricky at first. I realized that SVM isn't just about drawing a line to separate data. It's about finding the best way to divide things, even in more complicated spaces. I didn't have to dive into too many outside resources, but revisiting the basics of SVM helped me understand better.

Challenges and Triumphs. Training took a while, but once I got it working on the subset, it felt great. Seeing the model classify images reasonably well was a win.

Personal Growth. This exercise gave me a new appreciation for classical machine learning techniques. I also want to learn other image recognition methods because I want to focus my career on computer vision.

Looking Ahead. After finishing this exercise, I'm still curious about how SVM compares to deep learning models on more complex tasks. I'm excited to dive deeper into neural networks and learn more about models specifically built for handling images. This experience gave me a solid foundation, and I'm looking forward to building on it in future projects.