

▼ I. Project Overview



In this project, we'll build a neural network classifier that determines: **MUFFIN... or CHIHUAHUA!**

This is what we'll cover in the tutorial:

- 1) Build the neural network
- 2) Load the data
- 3) Train the model on the data
- 4) Visualize the results

Remember: This is an **INTERACTIVE** Notebook!

You should run and play with the code as you go to see how it works. Select a cell and **press shift-enter to execute code.**

The screenshot shows a Jupyter Notebook interface. At the top, there's a toolbar with various icons for file operations like saving, opening, and running cells. The title bar says "jupyter Untitled (unsaved changes)". On the right, there's a Python 2 logo. Below the toolbar, the main area has a heading "Simple Jupyter demo". A text block follows: "This cell has text formatted using the markdown language, which gets rendered like regular html. The next cell has some code:". A code cell is shown with the input:

```
In [57]: import random  
for i in range(3):  
    print random.random()  
x = 10
```

The output of this code is three floating-point numbers: 0.10564822904, 0.153941700348, and 0.518503128416. Below this, another text cell contains the text "Here is another text cell, with some *formatting*.":

▼ II. Deep Learning Tutorial

Let's get to the fun stuff!



Generic Python imports (select the below cell and press shift-enter to execute it)

```
import matplotlib.pyplot as plt # graphical library, to plot images  
# special Jupyter notebook command to show plots inline instead of in a new window  
%matplotlib inline
```

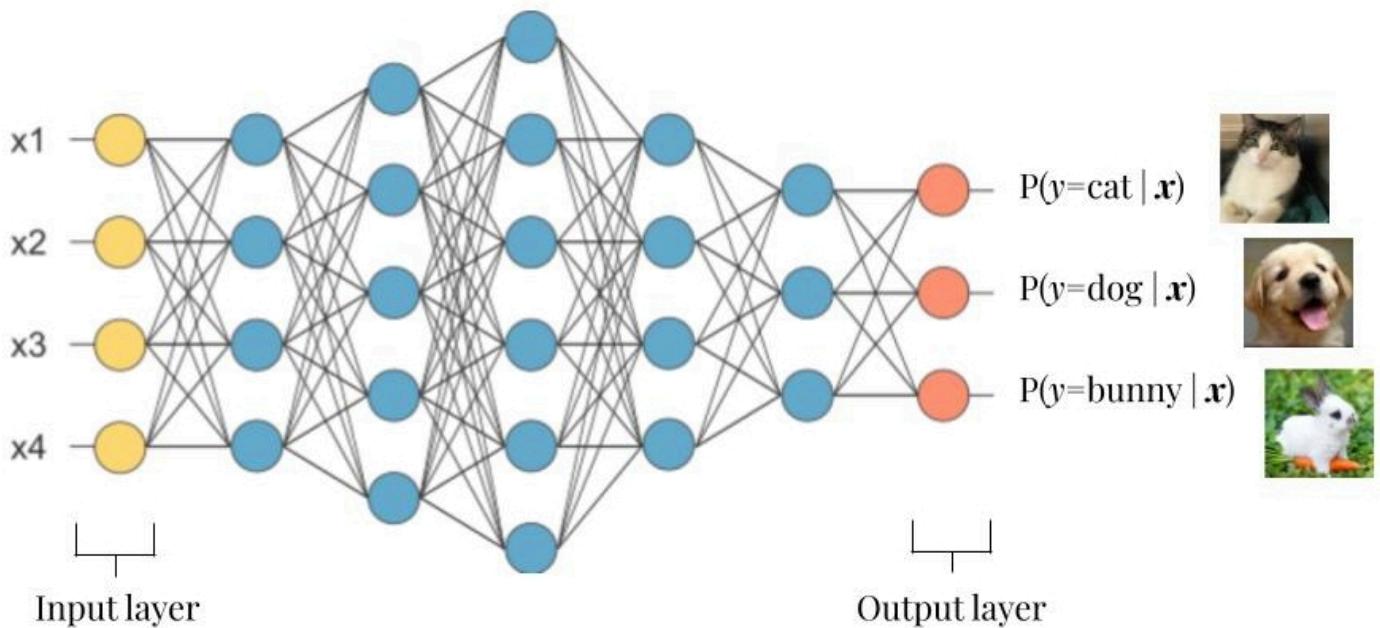
Deep learning imports

```
import torch
from torchvision import datasets, models, transforms
import torch.nn as nn
from torch.nn import functional as F
import torch.optim as optim
```

```
# PyTorch deep learning framework
# extension to PyTorch for distributed learning
# neural networks module of PyTorch
# special functions
# optimizers
```

▼ (1) Build our Neural Network

Recall from the lesson that a neural network generally looks like this. Input is on the left, output is on the right. The number of output neurons correspond to the number of classes.



So let's define a similar architecture for our 2-class muffin-vs-chihuahua classifier:

```
#define image height and width
input_height = 224
input_width = 224

# Extends PyTorch's neural network baseclass
class MySkynet(nn.Module):
    """
    A very basic neural network.
    """

    def __init__(self, input_dim=(3, input_height, input_width)):
        """
        Constructs a neural network.

        input_dim: a tuple that represents "channel x height x width" dimensions of
        """
        super().__init__()
        # the total number of RGB pixels in an image is the tensor's volume
        num_in_features = input_dim[0] * input_dim[1] * input_dim[2]
        # input layer
        self.layer_0 = nn.Linear(num_in_features, 128)
        # hidden layers
        self.layer_1 = nn.Linear(128, 64)
        self.layer_2 = nn.Linear(64, 32)
        # output layer, output size of 2 for chihuahua and muffin
        self.layer_3 = nn.Linear(32, 2)

    def forward(self, x):
        """
        Define the forward pass through our network.
        """

        batch_size = x.shape[0]
        # convert our RGB tensor into one long vector
        x = x.view(batch_size, -1)

        # pass through our layers
        x = F.relu(self.layer_0(x))
        x = F.relu(self.layer_1(x))
        x = F.relu(self.layer_2(x))
        x = F.relu(self.layer_3(x))

        # convert the raw output to probability predictions
        x = F.softmax(x, dim=1)

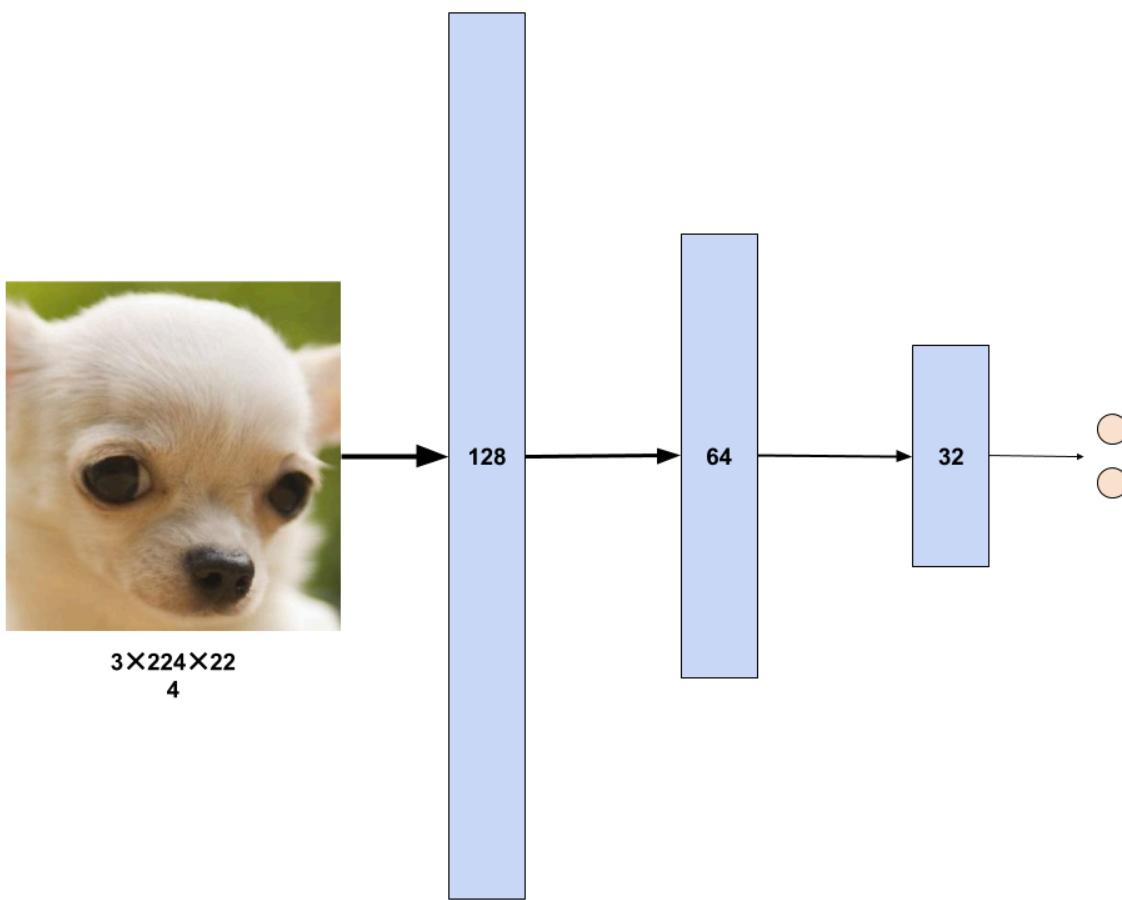
    return x
```

Now that we've defined the network above, let's initialize it. If available, we'll place the network on the GPU; if not, it goes on the CPU.

```
# cuda:0 means the first cuda device found
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model = MySkynet(input_dim=(3, input_height, input_width)).to(device)
model

→ MySkynet(
    (layer_0): Linear(in_features=150528, out_features=128, bias=True)
    (layer_1): Linear(in_features=128, out_features=64, bias=True)
    (layer_2): Linear(in_features=64, out_features=32, bias=True)
    (layer_3): Linear(in_features=32, out_features=2, bias=True)
)
```

Essentially, our network looks like this:



▼ (2) Data and Data Loading

Separate "train" and "test" datasets

Recall from the below slide, we should make two separate datasets to train and test our model. That way, we know our model learns more than rote memorization.

When is Your ML Model Ready?

Without care, ML models “memorize” answers from training data.

Solution: evaluate performance on **test** set, separate from the **train** set



Homework == Train Set



Test == ... Test Set

▼ Inspect our data

Let's look in our data folder to see what's there. As you can see, the folder is **split into "train" for training, and "validation" for testing** (to validate our model).

```
import os # interact with the os. in our case, we want to view the file system

print("Data contents:", os.listdir("data"))
print("Train contents:", os.listdir("data/train"))
print("Validation contents:", os.listdir("data/validation"))

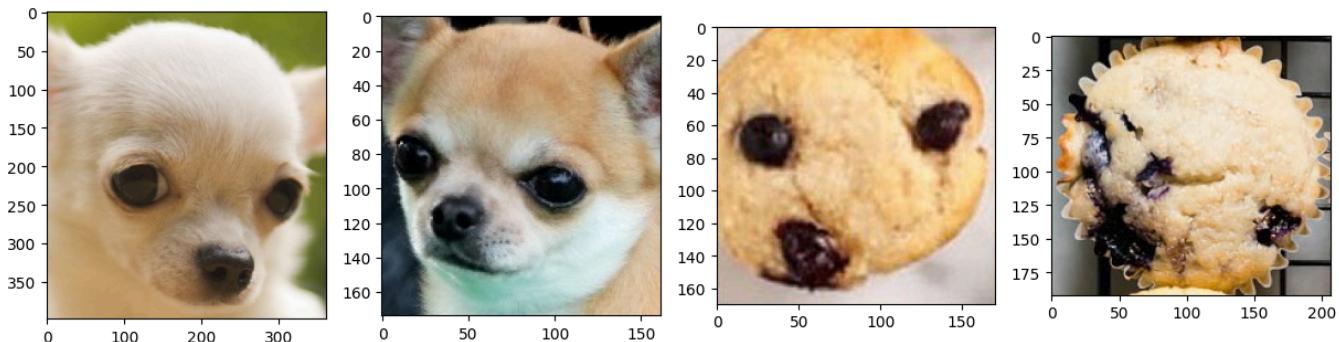
→ Data contents: ['train', 'validation']
  Train contents: ['chihuahua', 'muffin']
  Validation contents: ['chihuahua', 'muffin']
```

Let's also look at some of the images:

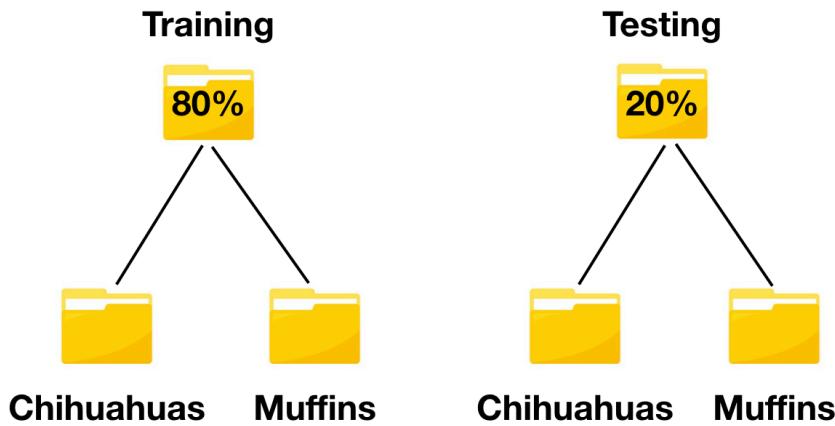
```
from PIL import Image # import our image opening tool

_, ax = plt.subplots(1, 4, figsize=(15,60)) # to show 4 images side by side, make a
ax[0].imshow(Image.open("data/train/chihuahua/4.jpg")) # show the chihuahua in the
ax[1].imshow(Image.open("data/train/chihuahua/5.jpg")) # show the chihuahua in the
ax[2].imshow(Image.open("data/train/muffin/131.jpg")) # show the muffin in the thi
ax[3].imshow(Image.open("data/train/muffin/107.jpg")) # show the muffin in the fou
```

→ <matplotlib.image.AxesImage at 0x789686569d50>



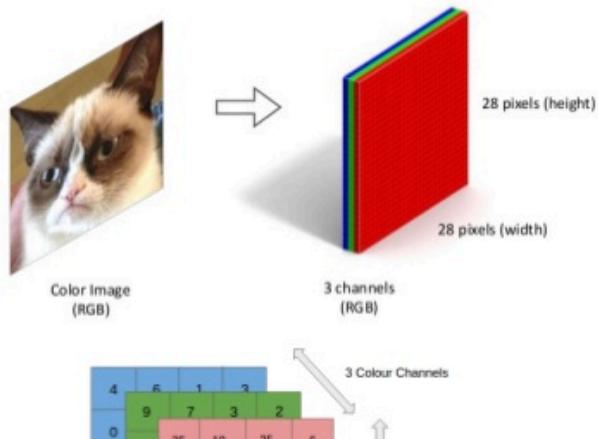
If you look in the data folder on your computer, there are 120 train images and 30 validation. So our data is split like this:



▼ Load our data

That's great that we have data! But we have to load all the images and convert them into a form that our neural network understands. Specifically, PyTorch works with **Tensor** objects. (A tensor is just a multidimensional matrix, i.e. an N-d array.)

color image is 3rd-order tensor



To easily convert our image data into tensors, we use the help of a "dataloader." The dataloader packages data into convenient boxes for our model to use. You can think of it like one person



passing boxes (tensors) to another.

First, we define some "transforms" to convert images to tensors. We must do so for both our train and validation datasets.

For more information about transforms, check out the link here:

<https://pytorch.org/docs/stable/torchvision/transforms.html>

```

normalize = transforms.Normalize(mean=[0.5, 0.5, 0.5],
                                std=[0.5, 0.5, 0.5])

# transforms for our training data
train_transforms = transforms.Compose([
    # resize to resnet input size
    transforms.Resize((input_height, input_width)),
    # transform image to PyTorch tensor object
    transforms.ToTensor(),
    normalize
])

# these validation transforms are exactly the same as our train transforms
validation_transforms = transforms.Compose([
    transforms.Resize((input_height, input_width)),
    transforms.ToTensor(),
    normalize
])

print("Train transforms:", train_transforms)

```

→ Train transforms: Compose(
 Resize(size=(224, 224), interpolation=bilinear, max_size=None, antialias=True)
 ToTensor()
 Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])
)

Second, we create the datasets, by passing the transforms into the ImageFolder constructor.

These just represent the folders that hold the images.

```

# insert respective transforms to replace ?
image_datasets = {
    'train':
        datasets.ImageFolder('data/train', train_transforms),
    'validation':
        datasets.ImageFolder('data/validation', validation_transforms)

print("==Train Dataset==\n", image_datasets["train"])
print()
print("==Validation Dataset==\n", image_datasets["train"])

→ ==Train Dataset==
Dataset ImageFolder
    Number of datapoints: 120
    Root location: data/train
    StandardTransform
    Transform: Compose(
        Resize(size=(224, 224), interpolation=bilinear, max_size=None, antialias=True)
        ToTensor()
    )

```

```

        Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])
    )

==Validation Dataset==
Dataset ImageFolder
    Number of datapoints: 120
    Root location: data/train
    StandardTransform
    Transform: Compose(
        Resize(size=(224, 224), interpolation=bilinear, max_size=None, an
              ToTensor()
        Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])
    )
)

```

And finally, form dataloaders from the datasets:

```

# define batch size, number of images to load in at once

dataloaders = {
    'train':
        torch.utils.data.DataLoader(
            image_datasets['train'],
            batch_size=4,
            shuffle=True,
            num_workers=4),
    'validation':
        torch.utils.data.DataLoader(
            image_datasets['validation'],
            batch_size=4,
            shuffle=False,
            num_workers=4)}

print("Train loader:", dataloaders["train"])
print("Validation loader:", dataloaders["validation"])

```

→ Train loader: <torch.utils.data.dataloader.DataLoader object at 0x78967d6363b0>
 Validation loader: <torch.utils.data.dataloader.DataLoader object at 0x78974a13c
 /usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: User
 warnings.warn(_create_warning_msg)

We can see a dataloader outputs 2 things: a BIG tensor to represent an image, and a vector to represent the labels (0 or 1).

```
next(iter(dataloaders["train"]))
```

→

```

[[ 0.9922,  0.9922,  0.9922,  ...,  0.9922,  0.9922,  0.9922],
 [ 0.9922,  0.9922,  0.9922,  ...,  0.9922,  0.9922,  0.9922],  

 ...,  

 [-0.7333, -0.6941, -0.7020,  ..., -0.0353, -0.0510, -0.1216],  

 [-0.5608, -0.5216, -0.5216,  ...,  0.0196,  0.0275, -0.0275],  

 [-0.3804, -0.3255, -0.3098,  ...,  0.0980,  0.1137,  0.0667]]],  

[[[ 0.6392,  0.6471,  0.6392,  ...,  0.3490,  0.3412,  0.3333],
 [ 0.6392,  0.6471,  0.6392,  ...,  0.3412,  0.3333,  0.3333],
 [ 0.6392,  0.6471,  0.6392,  ...,  0.3333,  0.3255,  0.3255],  

 ...,  

 [ 0.4588,  0.4431,  0.4510,  ...,  0.5216,  0.5294,  0.5216],
 [ 0.4902,  0.4902,  0.4824,  ...,  0.5137,  0.5137,  0.5137],
 [ 0.4902,  0.4902,  0.4902,  ...,  0.5059,  0.5059,  0.5137]],  

[[ 0.4667,  0.4745,  0.4667,  ...,  0.1294,  0.1216,  0.1216],
 [ 0.4667,  0.4745,  0.4667,  ...,  0.1216,  0.1137,  0.1216],
 [ 0.4588,  0.4667,  0.4588,  ...,  0.1137,  0.0980,  0.1137],  

 ...,  

 [ 0.3882,  0.3961,  0.3882,  ...,  0.3176,  0.3255,  0.3176],
 [ 0.4196,  0.4196,  0.4118,  ...,  0.3098,  0.3098,  0.3098],
 [ 0.4118,  0.4118,  0.4118,  ...,  0.3020,  0.3020,  0.3098]],  

[[ 0.3804,  0.3882,  0.3804,  ..., -0.0431, -0.0510, -0.0431],
 [ 0.3804,  0.3882,  0.3804,  ..., -0.0431, -0.0510, -0.0431],
 [ 0.3961,  0.4039,  0.3961,  ..., -0.0431, -0.0431, -0.0353],  

 ...,  

 [ 0.4118,  0.4118,  0.4039,  ...,  0.0588,  0.0667,  0.0588],
 [ 0.4431,  0.4431,  0.4353,  ...,  0.0510,  0.0510,  0.0510],
 [ 0.4196,  0.4196,  0.4196,  ...,  0.0431,  0.0431,  0.0510]]],  

[[[-0.5137, -0.4824, -0.4510,  ..., -0.4510, -0.5216, -0.4588],
 [-0.4980, -0.4824, -0.4431,  ..., -0.4431, -0.3569, -0.2078],
 [-0.4824, -0.4745, -0.4196,  ..., -0.4118, -0.3333, -0.2314],  

 ...,  

 [-0.3569, -0.3725, -0.4196,  ..., -0.4353, -0.4588, -0.4902],
 [-0.1765, -0.2235, -0.2941,  ..., -0.4039, -0.4275, -0.4510],
 [-0.0510, -0.0824, -0.1059,  ..., -0.4118, -0.4196, -0.4275]],  

[[[-0.5608, -0.5294, -0.4980,  ..., -0.5451, -0.6078, -0.5373],
 [-0.5451, -0.5294, -0.4902,  ..., -0.5373, -0.4431, -0.2784],
 [-0.5294, -0.5216, -0.4667,  ..., -0.5137, -0.4196, -0.3020],  

 ...,  

 [-0.2392, -0.2706, -0.3176,  ..., -0.4196, -0.4431, -0.4745],
 [-0.0431, -0.0980, -0.1686,  ..., -0.3882, -0.4196, -0.4353],
 [ 0.0902,  0.0588,  0.0196,  ..., -0.3961, -0.4118, -0.4275]],  

[[[-0.6863, -0.6549, -0.6157,  ..., -0.6706, -0.7255, -0.6314],
 [-0.6706, -0.6549, -0.6078,  ..., -0.6627, -0.5608, -0.3961],
 [-0.6706, -0.6471, -0.5843,  ..., -0.6392, -0.5451, -0.4196],  

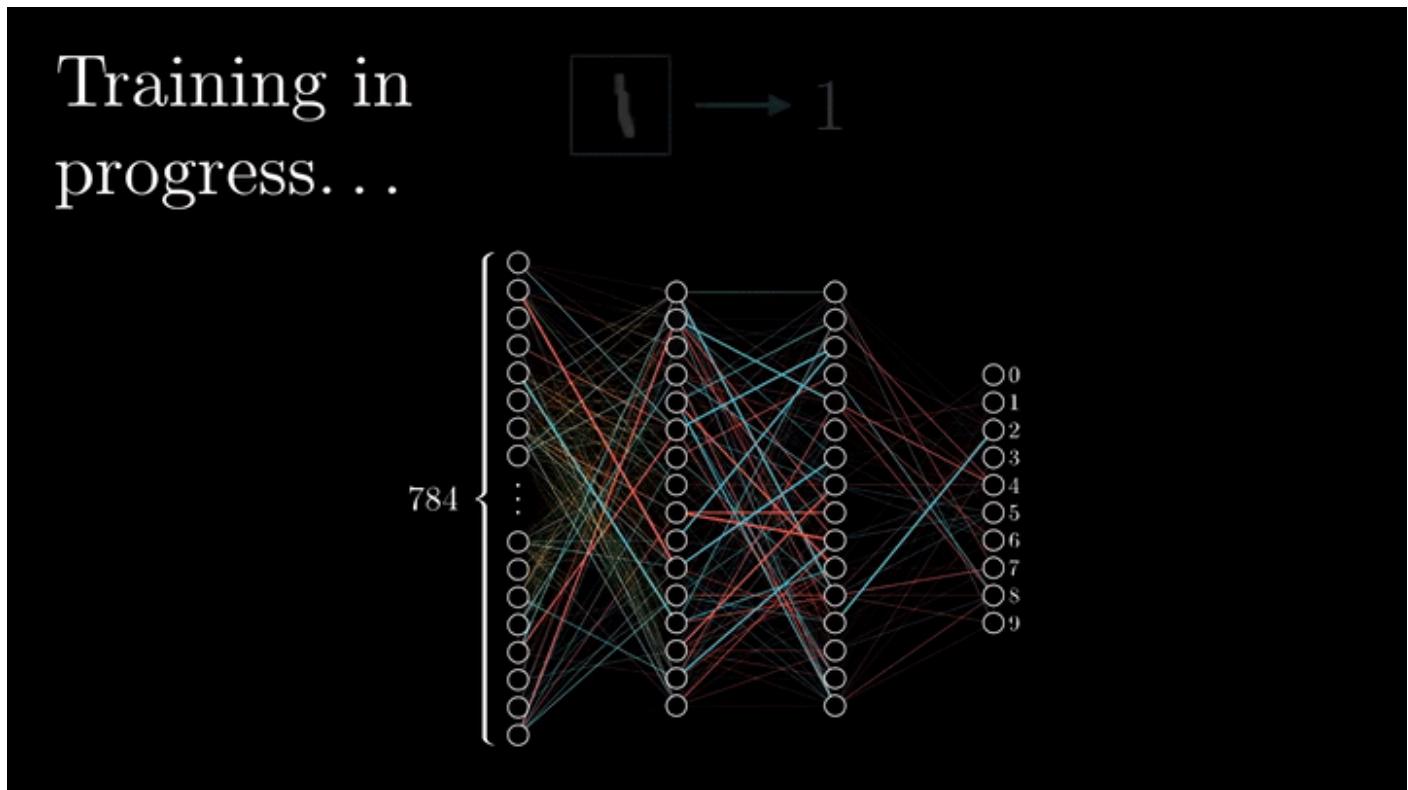
 ...,  

 [-0.1050, -0.1216, -0.1204,  ..., -0.5216, -0.5451, -0.5765]
]
```

❖ (4) Train the model!

Hurray! We've built a neural network and have data to give it. Now we **repeatedly iterate over the data to train the model**.

Every time the network gets a new example, it looks something like this. Note the **forward pass** and the corresponding **backward pass**.



❖ Define the train loop

We want the network to learn from every example in our training dataset. However, the best performance comes from more practice. Therefore, we **run through our dataset for multiple epochs**.

After each epoch, we'll check how our model performs on the validation set to monitor its progress.

```
from tqdm import trange, tqdm_notebook # import progress bars to show train progress

def train_model(model, dataloaders, loss_function, optimizer, num_epochs):
    """
    Trains a model using the given loss function and optimizer, for a certain number

    model: a PyTorch neural network
    loss_function: a mathematical function that compares predictions and labels to r
    num_epochs: the number of times to run through the full training dataset
    """

    # train for n epochs. an epoch is a full iteration through our dataset
    for epoch in trange(num_epochs, desc="Total progress", unit="epoch"):
        # print a header
        print('Epoch {} / {}'.format(epoch+1, num_epochs))
        print('-----')

        # first train over the dataset and update weights; at the end, calculate our
        for phase in ['train', 'validation']:
            if phase == 'train':
                model.train()
            else:
                model.eval()

            # keep track of the overall loss and accuracy for this batch
            running_loss = 0.0
            running_corrects = 0

            # iterate through the inputs and labels in our dataloader
            # (the tqdm_notebook part is to display a progress bar)
            for inputs, labels in tqdm_notebook(dataloaders[phase], desc=phase, unit
                # move inputs and labels to appropriate device (GPU or CPU)
                inputs = inputs.to(device)
                labels = labels.to(device)

                # FORWARD PASS
                outputs = model(inputs)
                # compute the error of the model's predictions
                loss = loss_function(outputs, labels)

                if phase == 'train':
                    # BACKWARD PASS
                    optimizer.zero_grad() # clear the previous gradients
                    loss.backward() # backpropagate the current error gradient
                    optimizer.step() # update the weights (i.e. do the learning)

                    # track our accumulated loss
                    running_loss += loss.item() * inputs.size(0)
                    # track number of correct to compute accuracy
                    _, preds = torch.max(outputs, 1)
                    running_corrects += torch.sum(preds == labels.data)

    # track our accumulated loss
    running_loss /= len(dataloaders['train'].dataset)
    # track number of correct to compute accuracy
    running_corrects /= len(dataloaders['train'].dataset)
    accuracy = running_corrects / len(dataloaders['train'].dataset)

    return accuracy
```

```
# print our progress
epoch_loss = running_loss / len(image_datasets[phase])
epoch_acc = running_corrects.double() / len(image_datasets[phase])
print(f'{phase} error: {epoch_loss:.4f}, Accuracy: {epoch_acc:.4f}')

print()
```

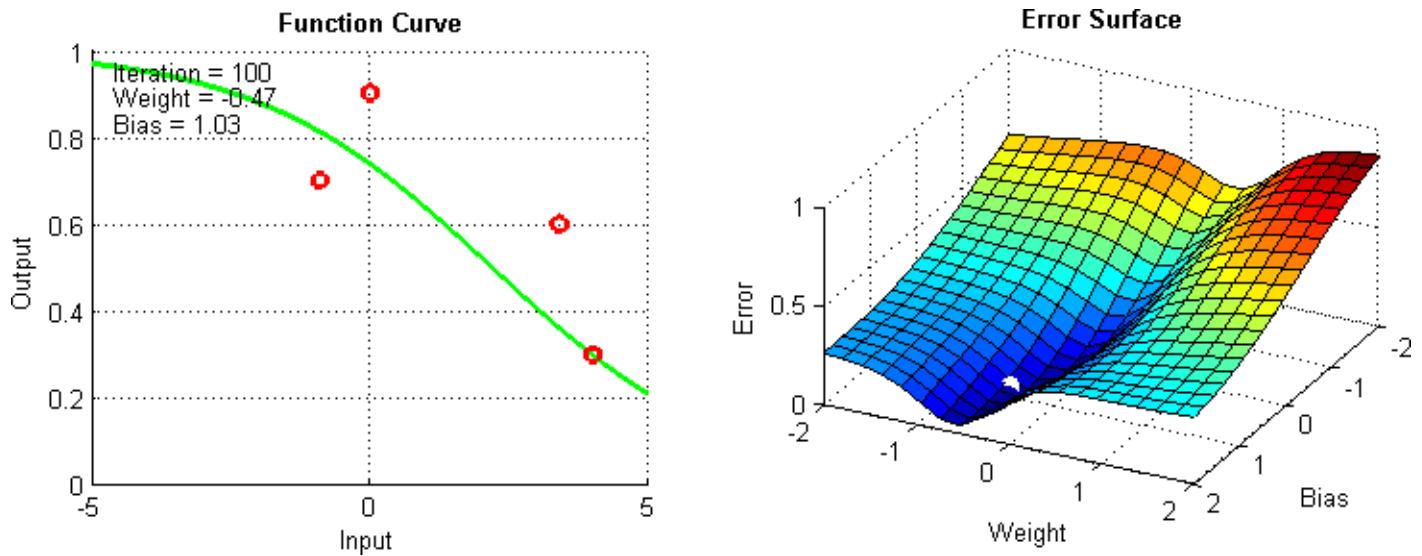
▼ Loss function and optimizer

One last thing: we must define a function that gives feedback for how well the model performs. This is the **loss**, or "error" **function**, that compares model predictions to the true labels.

Once we calculate the error, we also need to define how the model should react to that feedback.

The optimizer determines how the network learns from feedback.

```
loss_function = nn.CrossEntropyLoss() # the most common error function
optimizer = optim.SGD(model.parameters(), lr=0.1) # Stochastic Gradient Descent, wi
```



▼ Run training

Let's put everything together and TRAIN OUR MODEL! =D

```
train_model(model, dataloaders, loss_function, optimizer, num_epochs=3)
```

```
→ <ipython-input-22-c9fe14cbe6c0>:12: TqdmDeprecationWarning: Please use `tqdm.not  
for epoch in tnrangne(num_epochs, desc="Total progress", unit="epoch"):  
    Total progress: 100%                                         3/3 [00:15<00:00, 4.96s/epoch]  
Epoch 1/3  
-----  
<ipython-input-22-c9fe14cbe6c0>:30: TqdmDeprecationWarning: This function will b  
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`  
    for inputs, labels in tqdm_notebook(dataloaders[phase], desc=phase, unit="batch"):  
        train error: 0.6366, Accuracy: 0.7167  
        validation error: 0.5204, Accuracy: 0.8000  
  
Epoch 2/3  
-----  
    train error: 0.5664, Accuracy: 0.7333  
    validation error: 0.4762, Accuracy: 0.8333  
  
Epoch 3/3  
-----  
    train error: 0.4282, Accuracy: 0.8917  
    validation error: 0.4738, Accuracy: 0.8333
```

▼ Examine model performance



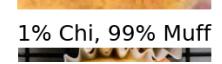
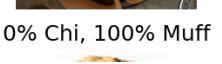
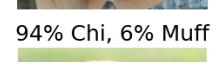
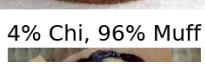
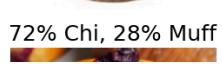
How do we examine our model's predictions? Let's visualize what the model thinks on the validation set.

```
from glob import glob
from math import floor

# get all the images from our validation sets
validation_img_paths = glob("data/validation/**/*.jpg", recursive=True)
images = [Image.open(img_path) for img_path in validation_img_paths]

# put all the images together to run through our model
validation_batch = torch.stack( [validation_transforms(img).to(device) for img in images])
pred_logits_tensor = model(validation_batch)
pred_probs = pred_logits_tensor.cpu().data.numpy()

# show the probabilities for each picture
fig, axs = plt.subplots(6, 5, figsize=(20, 20))
for i, img in enumerate(images):
    ax = axs[floor(i/5)][i % 5]
    ax.axis('off')
    ax.set_title("{:.0f}% Chi, {:.0f}% Muff".format(100*pred_probs[i,0], 100*pred_probs[i,1]))
    ax.imshow(img)
```



Consider: How accurate was your model? How confident were its predictions? Does it make clear-